

# Feature\_Engineering (1)

August 19, 2022

```
[1]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from pickle import dump
```

```
[2]: data = pd.DataFrame({'ID' : [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
                          'Name' : ['John', 'Jack', 'Mariah', 'Krishna', 'Danny', 'Lisa',
↳ 'Andrew', 'Ravi', 'Garima', 'Kavita'],
                          'Gender' : ['Male', 'Male', 'Female', 'Male', 'Male', 'Female',
↳ 'Male', 'Male', 'Female', 'Female'],
                          'Profession' : ['Manager', 'Manager', 'Developer', 'Team Lead',
↳ 'Team Lead', 'Assistant Manager', 'Assistant Manager', 'Manager', 'Assistant
↳ Manager', 'CEO'],
                          'Salary' : [100000, 120000, 95000, 99000, 105000, 145000, 155000,
↳ 78000, 167000, 195000],
                          'Experience' : [10, 13, 6, 8, 10, 13, 19, 15, 20, 24]})
```

```
[3]: data
```

```
[3]:
```

	ID	Name	Gender	Profession	Salary	Experience
0	1	John	Male	Manager	100000	10
1	2	Jack	Male	Manager	120000	13
2	3	Mariah	Female	Developer	95000	6
3	4	Krishna	Male	Team Lead	99000	8
4	5	Danny	Male	Team Lead	105000	10
5	6	Lisa	Female	Assistant Manager	145000	13
6	7	Andrew	Male	Assistant Manager	155000	19
7	8	Ravi	Male	Manager	78000	15
8	9	Garima	Female	Assistant Manager	167000	20
9	10	Kavita	Female	CEO	195000	24

```
[4]: data.describe()
```

```
[4]:
```

	ID	Salary	Experience
count	10.00000	10.000000	10.000000
mean	5.50000	125900.000000	13.800000
std	3.02765	37698.953714	5.731007

min	1.00000	78000.000000	6.000000
25%	3.25000	99250.000000	10.000000
50%	5.50000	112500.000000	13.000000
75%	7.75000	152500.000000	18.000000
max	10.00000	195000.000000	24.000000

### 0.0.1 Explore the data types

```
[5]: data.dtypes
```

```
[5]: ID          int64
     Name        object
     Gender       object
     Profession   object
     Salary       int64
     Experience   int64
     dtype: object
```

### 0.0.2 Separating categorical and numeric features

```
[59]: dictionary={}
     dictionary['num'] = data.dtypes[data.dtypes=='int64'].index
     dictionary['cat'] = data.dtypes[data.dtypes=='object'].index
     dictionary
```

```
[59]: {'num': Index(['ID', 'Salary', 'Experience'], dtype='object'),
      'cat': Index(['Name', 'Gender', 'Profession'], dtype='object')}
```

### 0.0.3 Cardinality of Categorical Variables

```
[60]: data['Name'].unique()
```

```
[60]: array(['John', 'Jack', 'Mariah', 'Krishna', 'Danny', 'Lisa', 'Andrew',
      'Ravi', 'Garima', 'Kavita'], dtype=object)
```

```
[61]: data['Gender'].unique()
```

```
[61]: array(['Male', 'Female'], dtype=object)
```

```
[62]: data['Profession'].unique()
```

```
[62]: array(['Manager', 'Developer', 'Team Lead', 'Assistant Manager', 'CEO'],
      dtype=object)
```

#### 0.0.4 Frequency of categories in a categorical variable

```
[63]: data['Name'].value_counts()
```

```
[63]: Danny      1
      Krishna    1
      Kavita     1
      John       1
      Ravi       1
      Mariah     1
      Garima     1
      Jack       1
      Lisa       1
      Andrew     1
      Name: Name, dtype: int64
```

```
[64]: data['Gender'].value_counts()
```

```
[64]: Male        6
      Female      4
      Name: Gender, dtype: int64
```

```
[65]: data['Profession'].value_counts()
```

```
[65]: Manager          3
      Assistant Manager 3
      Team Lead        2
      Developer         1
      CEO              1
      Name: Profession, dtype: int64
```


```
[66]: data['Gender'].hist()
```

```
[66]: <AxesSubplot:>
```




```
[67]: data['Profession'].hist()
```

```
[67]: <AxesSubplot:>
```




```
[68]: sns.countplot(x='Gender',data=data,palette='hls')  
plt.show()
```



### Numerical Variables

```
[69]: data['Salary'].hist(bins=5)
```

```
[69]: <AxesSubplot:>
```



```
[70]: data[dictionary['cat']]
```

```
[70]:
```

	Name	Gender	Profession
0	John	Male	Manager
1	Jack	Male	Manager
2	Mariah	Female	Developer
3	Krishna	Male	Team Lead
4	Danny	Male	Team Lead
5	Lisa	Female	Assistant Manager
6	Andrew	Male	Assistant Manager
7	Ravi	Male	Manager
8	Garima	Female	Assistant Manager
9	Kavita	Female	CEO

```
[71]: from sklearn.preprocessing import MinMaxScaler, StandardScaler, LabelEncoder
```

```
[72]: # Profession column is ordinal data. Hence I can use the label encoder
```

```
le_profession = LabelEncoder()
le_profession.fit(data['Profession'])
```

```
[72]: LabelEncoder()
```

```
[73]: le_profession.classes_
```

```
[73]: array(['Assistant Manager', 'CEO', 'Developer', 'Manager', 'Team Lead'],
      dtype=object)
```

fit method learns the transformer's attributes. It does not change the column. To change the column we will now use the learnt attribute to do the transformation by using transform function

```
[74]: data['Profession'] = le_profession.transform(data['Profession'])
data
```

```
[74]:
```

	ID	Name	Gender	Profession	Salary	Experience
0	1	John	Male	3	100000	10
1	2	Jack	Male	3	120000	13
2	3	Mariah	Female	2	95000	6
3	4	Krishna	Male	4	99000	8
4	5	Danny	Male	4	105000	10
5	6	Lisa	Female	0	145000	13
6	7	Andrew	Male	0	155000	19
7	8	Ravi	Male	3	78000	15
8	9	Garima	Female	0	167000	20
9	10	Kavita	Female	1	195000	24

```
[75]: le_profession.inverse_transform(data['Profession'])
```

```
[75]: array(['Manager', 'Manager', 'Developer', 'Team Lead', 'Team Lead',
      'Assistant Manager', 'Assistant Manager', 'Manager',
      'Assistant Manager', 'CEO'], dtype=object)
```

```
[76]: # Gender is nominal data because it cannot be ranked. Hence we will apply one-
      ↪hot encoder

pd.get_dummies(data['Gender'])
```

```
[76]:
```

	Female	Male
0	0	1
1	0	1
2	1	0
3	0	1
4	0	1
5	1	0
6	0	1
7	0	1
8	1	0
9	1	0

```
[77]: data.drop(['Name'], inplace = True, axis = 1)
data
```

```
[77]:
```

	ID	Gender	Profession	Salary	Experience
0	1	Male	3	100000	10
1	2	Male	3	120000	13
2	3	Female	2	95000	6
3	4	Male	4	99000	8
4	5	Male	4	105000	10
5	6	Female	0	145000	13
6	7	Male	0	155000	19
7	8	Male	3	78000	15
8	9	Female	0	167000	20
9	10	Female	1	195000	24

```
[78]: data = pd.get_dummies(data)
data
```

```
[78]:
```

	ID	Profession	Salary	Experience	Gender_Female	Gender_Male
0	1	3	100000	10	0	1
1	2	3	120000	13	0	1
2	3	2	95000	6	1	0
3	4	4	99000	8	0	1
4	5	4	105000	10	0	1
5	6	0	145000	13	1	0
6	7	0	155000	19	0	1
7	8	3	78000	15	0	1
8	9	0	167000	20	1	0
9	10	1	195000	24	1	0

### Scaling Numerical Variables using MinMax

```
[80]: min_max_sal = MinMaxScaler()
min_max_sal.fit(data[['Salary']])
```

```
[80]: MinMaxScaler()
```

```
[82]: min_max_sal.data_max_, min_max_sal.data_min_
```

```
[82]: (array([195000.]), array([78000.]))
```

```
[83]: data
```

```
[83]:
```

	ID	Profession	Salary	Experience	Gender_Female	Gender_Male
0	1	3	100000	10	0	1
1	2	3	120000	13	0	1
2	3	2	95000	6	1	0
3	4	4	99000	8	0	1
4	5	4	105000	10	0	1
5	6	0	145000	13	1	0



6	7	0	155000	19	0	1
7	8	3	78000	15	0	1
8	9	0	167000	20	1	0
9	10	1	195000	24	1	0

```
[16]: min_max_sal.data_max_
```

```
[16]: array([195000.])
```

```
[84]: data['Salary_min_max_scaled'] = min_max_sal.transform(data[['Salary']])
data
```

```
[84]:
```

	ID	Profession	Salary	Experience	Gender_Female	Gender_Male	\
0	1	3	100000	10	0	1	
1	2	3	120000	13	0	1	
2	3	2	95000	6	1	0	
3	4	4	99000	8	0	1	
4	5	4	105000	10	0	1	
5	6	0	145000	13	1	0	
6	7	0	155000	19	0	1	
7	8	3	78000	15	0	1	
8	9	0	167000	20	1	0	
9	10	1	195000	24	1	0	

	Salary_min_max_scaled
0	0.188034
1	0.358974
2	0.145299
3	0.179487
4	0.230769
5	0.572650
6	0.658120
7	0.000000
8	0.760684
9	1.000000

```
[85]: data.describe()
```

```
[85]:
```

	ID	Profession	Salary	Experience	Gender_Female	\
count	10.00000	10.000000	10.000000	10.000000	10.000000	
mean	5.50000	2.000000	125900.000000	13.800000	0.400000	
std	3.02765	1.632993	37698.953714	5.731007	0.516398	
min	1.00000	0.000000	78000.000000	6.000000	0.000000	
25%	3.25000	0.250000	99250.000000	10.000000	0.000000	
50%	5.50000	2.500000	112500.000000	13.000000	0.000000	
75%	7.75000	3.000000	152500.000000	18.000000	1.000000	
max	10.00000	4.000000	195000.000000	24.000000	1.000000	

	Gender_Male	Salary_min_max_scaled
count	10.000000	10.000000
mean	0.600000	0.409402
std	0.516398	0.322213
min	0.000000	0.000000
25%	0.000000	0.181624
50%	1.000000	0.294872
75%	1.000000	0.636752
max	1.000000	1.000000

### Scaling Numerical Variables using StandardScaler

```
[86]: std_scale_sal = StandardScaler()
std_scale_sal.fit(data[['Salary']])
```

```
[86]: StandardScaler()
```

```
[87]: std_scale_sal.mean_, std_scale_sal.scale_
```

```
[87]: (array([125900.]), array([35764.36774221]))
```

```
[88]: data['Salary_standard_scaled'] = std_scale_sal.transform(data[['Salary']])
data
```

```
[88]:
```

	ID	Profession	Salary	Experience	Gender_Female	Gender_Male	\
0	1	3	100000	10	0	1	
1	2	3	120000	13	0	1	
2	3	2	95000	6	1	0	
3	4	4	99000	8	0	1	
4	5	4	105000	10	0	1	
5	6	0	145000	13	1	0	
6	7	0	155000	19	0	1	
7	8	3	78000	15	0	1	
8	9	0	167000	20	1	0	
9	10	1	195000	24	1	0	

	Salary_min_max_scaled	Salary_standard_scaled
0	0.188034	-0.724184
1	0.358974	-0.164969
2	0.145299	-0.863988
3	0.179487	-0.752145
4	0.230769	-0.584381
5	0.572650	0.534051
6	0.658120	0.813659
7	0.000000	-1.339322
8	0.760684	1.149188

9

1.000000

1.932091

```
[89]: data.describe()
```

```
[89]:
```

	ID	Profession	Salary	Experience	Gender_Female \
count	10.00000	10.000000	10.000000	10.000000	10.000000
mean	5.50000	2.000000	125900.000000	13.800000	0.400000
std	3.02765	1.632993	37698.953714	5.731007	0.516398
min	1.00000	0.000000	78000.000000	6.000000	0.000000
25%	3.25000	0.250000	99250.000000	10.000000	0.000000
50%	5.50000	2.500000	112500.000000	13.000000	0.000000
75%	7.75000	3.000000	152500.000000	18.000000	1.000000
max	10.00000	4.000000	195000.000000	24.000000	1.000000

	Gender_Male	Salary_min_max_scaled	Salary_standard_scaled
count	10.000000	10.000000	1.000000e+01
mean	0.600000	0.409402	-2.220446e-17
std	0.516398	0.322213	1.054093e+00
min	0.000000	0.000000	-1.339322e+00
25%	0.000000	0.181624	-7.451551e-01
50%	1.000000	0.294872	-3.746746e-01
75%	1.000000	0.636752	7.437570e-01
max	1.000000	1.000000	1.932091e+00

```
[22]: std_scale_sal.inverse_transform(data[['Salary_standard_scaled']])
```

```
[22]: array([[100000.],
        [120000.],
        [ 95000.],
        [ 99000.],
        [105000.],
        [145000.],
        [155000.],
        [ 78000.],
        [167000.],
        [195000.]])
```

```
[91]: dump(std_scale_sal,open('std_scaler.pkl','wb'))
```

```
[ ]:
```