

# PCA

August 19, 2022

```
[1]: from sklearn.datasets import load_digits
```

```
[2]: import pandas as pd
import numpy as np
```

```
[3]: data = load_digits()
```

```
[4]: print(data.DESCR)
```

```
.. _digits_dataset:
```

Optical recognition of handwritten digits dataset

-----

**\*\*Data Set Characteristics:\*\***

```
:Number of Instances: 1797
:Number of Attributes: 64
:Attribute Information: 8x8 image of integer pixels in the range 0..16.
:Missing Attribute Values: None
:Creator: E. Alpaydin (alpaydin '@' boun.edu.tr)
:Date: July; 1998
```

This is a copy of the test set of the UCI ML hand-written digits datasets  
<https://archive.ics.uci.edu/ml/datasets/Optical+Recognition+of+Handwritten+Digits>

The data set contains images of hand-written digits: 10 classes where each class refers to a digit.

Preprocessing programs made available by NIST were used to extract normalized bitmaps of handwritten digits from a preprinted form. From a total of 43 people, 30 contributed to the training set and different 13 to the test set. 32x32 bitmaps are divided into nonoverlapping blocks of 4x4 and the number of on pixels are counted in each block. This generates an input matrix of 8x8 where each element is an integer in the range 0..16. This reduces dimensionality and gives invariance to small distortions.

For info on NIST preprocessing routines, see M. D. Garriss, J. L. Blue, G. T. Candela, D. L. Dimmick, J. Geist, P. J. Grother, S. A. Janet, and C. L. Wilson, NIST Form-Based Handprint Recognition System, NISTIR 5469, 1994.

.. topic:: References

- C. Kaynak (1995) Methods of Combining Multiple Classifiers and Their Applications to Handwritten Digit Recognition, MSc Thesis, Institute of Graduate Studies in Science and Engineering, Bogazici University.
- E. Alpaydin, C. Kaynak (1998) Cascading Classifiers, Kybernetika.
- Ken Tang and Ponnuthurai N. Suganthan and Xi Yao and A. Kai Qin. Linear dimensionality reduction using relevance weighted LDA. School of Electrical and Electronic Engineering Nanyang Technological University. 2005.
- Claudio Gentile. A New Approximate Maximal Margin Classification Algorithm. NIPS. 2000.

```
[5]: features = data.data
     target = data.target

     features.shape, target.shape
```

```
[5]: ((1797, 64), (1797,))
```

```
[6]: from sklearn.model_selection import train_test_split
```

```
[7]: X_train, X_test, y_train, y_test = train_test_split(features, target,
     ↪train_size = 0.6, random_state = 3)
     print(X_train.shape)
     print(X_test.shape)
     print(y_train.shape)
     print(y_test.shape)
```

```
(1078, 64)
(719, 64)
(1078,)
(719,)
```

```
[8]: from sklearn.linear_model import LogisticRegression
```

```
[9]: my_model = LogisticRegression()
```

```
[10]: my_model.fit(X_train, y_train)
```

```
/usr/local/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:765:
```

ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG)

[10]: LogisticRegression()

[11]: preds = my\_model.predict(X\_test)

[12]: from sklearn.metrics import accuracy\_score, confusion\_matrix

[13]: print(accuracy\_score(y\_test, preds))

0.9582753824756607

[14]: print(confusion\_matrix(y\_test, preds))

```
[[74  0  0  0  1  0  0  0  0  0]
 [ 0 72  0  0  0  0  0  0  2  1]
 [ 0  1 77  0  0  0  0  0  0  0]
 [ 0  0  2 66  0  0  0  0  1  0]
 [ 0  4  0  0 79  0  1  1  1  0]
 [ 0  1  0  0  1 60  0  0  0  2]
 [ 0  1  0  0  0  1 55  0  1  0]
 [ 0  0  0  1  0  0  0 73  2  0]
 [ 0  1  0  0  0  1  0  0 66  0]
 [ 0  0  0  0  1  1  0  0  1 67]]
```

[15]: from sklearn.decomposition import PCA

Initially I have 64 dimensions. Now I want to bring it down to 4 dimensions

[61]: pca = PCA(n\_components=.99)

[62]: pca.fit(X\_train)

[62]: PCA(n\_components=0.99)

[63]: pca.explained\_variance\_ratio\_

[63]: array([0.14924092, 0.13473079, 0.1156578 , 0.08302537, 0.06011606,  
0.05032773, 0.04284684, 0.03556855, 0.03294894, 0.03039433,  
0.02454552, 0.02311806, 0.01912235, 0.01765833, 0.01483066,

```
0.01426726, 0.01361367, 0.01250578, 0.01025873, 0.00956485,  
0.00899525, 0.00826308, 0.00757479, 0.00733326, 0.00666861,  
0.00603645, 0.00573677, 0.00514727, 0.00493476, 0.00431694,  
0.00372314, 0.00357927, 0.00333257, 0.00320508, 0.00296291,  
0.00287043, 0.0026174 , 0.0023106 , 0.00222457, 0.00213944,  
0.00178448])
```

```
[64]: X_train_transformed = pca.transform(X_train)  
X_test_transformed = pca.transform(X_test)
```

```
[65]: X_train_transformed.shape, X_test_transformed.shape
```

```
[65]: ((1078, 41), (719, 41))
```

```
[66]: X_train_transformed
```

```
[66]: array([[ 8.31885532, -11.63536164,  1.93999083, ...,  0.70688983,  
          -0.28709199, -0.10932069],  
          [ 3.69591051, 22.68576943, -5.66316747, ...,  1.32126734,  
          -2.25843322,  0.22766461],  
          [ 8.7400552 , -2.00460958, 12.21941747, ...,  0.42945575,  
          -0.28278685, -1.00717996],  
          ...,  
          [-0.20762926, 20.34392093, -2.0983895 , ..., -1.18724655,  
          -0.36066818, -0.82807739],  
          [19.78689345, -6.56299576, 18.79910461, ..., -1.92655226,  
          -3.21873055, -2.71094689],  
          [-12.92312023, -11.63513571, 14.44423877, ..., -0.49248432,  
           0.07430153,  0.55638722]])
```

```
[67]: pca.inverse_transform(X_train_transformed).shape
```

```
[67]: (1078, 64)
```

```
[68]: my_model.fit(X_train_transformed,y_train)
```

```
/usr/local/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:765:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG)

```
[68]: LogisticRegression()
```

```
[69]: preds = my_model.predict(X_test_transformed)
```

```
[70]: print(accuracy_score(y_test, preds))
```

```
0.9485396383866481
```

```
[ ]:
```